



THE GRAINGER COLLEGE OF ENGINEERING

# CS 521

## Technological Foundations of Blockchain and Cryptocurrency

*Grigore Rosu*

Topic 4 – Ethereum

I ILLINOIS

# Bitcoin's Limitations – Why Something More?

I

- Lack of Turing-completeness
  - Script language supports no loops; simulating them requires duplicating code
- Value-blindness
  - UTXO is all-or-nothing; no fine-grained control over withdrawal amounts
- Lack of state
  - UTXO are either spent or unspent; no multi-stage contracts or persistent variables
- Blockchain-blindness
  - Scripts cannot access block data (nonce, timestamp, previous hash)



# Ethereum

A next-generation smart contract  
and decentralized application platform

# Ethereum: A Programmable Blockchain

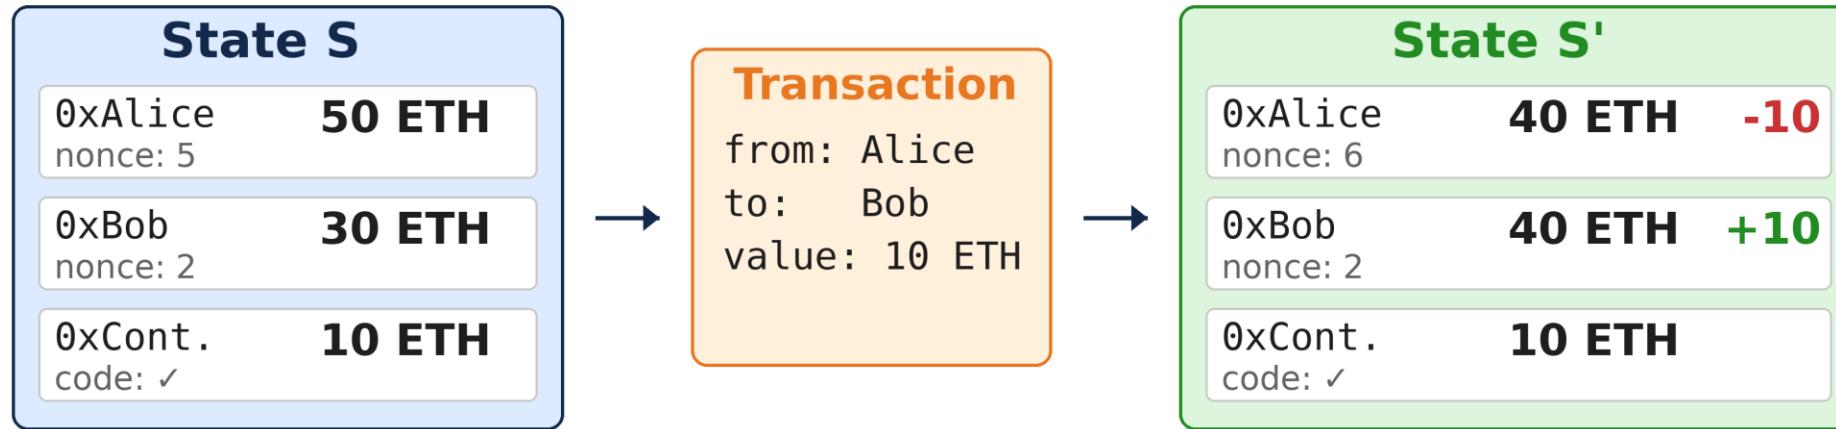
- Proposed by Vitalik Buterin in late 2013; whitepaper published 2014
- Goal: a blockchain with a built-in Turing-complete programming language
- Anyone can write smart contracts and decentralized applications (dApps)
  - Users define their own logic in a few lines of code
- Create arbitrary rules for ownership, transaction formats, and state transitions
- Not just digital cash – a general-purpose decentralized world computer
  - Tokens, DeFi, NFTs, DAOs, identity, governance, and more
- Launched July 30, 2015 (Frontier release)

# Ethereum as a State Transition System

- State: mapping from addresses (160-bit) to account states
  - Every account has: nonce, balance, storage root, code hash
- State transition function:  $\text{APPLY}(S, \text{TX}) \rightarrow S'$  or ERROR
  - Validates signature, deducts gas fees, runs code, refunds unused gas
- Transactions morph the state incrementally, block by block
- Key difference from Bitcoin:
  - Bitcoin = UTXO model (stateless); Ethereum = account model (stateful)
- Ethereum blocks contain both the transaction list AND the most recent state (root hash)
- State stored in a Modified Merkle Patricia Trie
  - Allows efficient state lookups and proofs without storing full history

# Ethereum State Transition Function

I



**APPLY( S , TX ) → S' or ERROR**

Yellow Paper:  $\sigma(t+1) \equiv Y(\sigma(t), T)$

APPLY({ Alice: 50, Bob: 30 }, "send 70")

→ **ERROR (insufficient balance: 50 < 70)**

# Ethereum Accounts

- Two types of accounts sharing the same 20-byte address space:
  - **Externally Owned Accounts (EOAs)**
    - Controlled by private keys (humans/wallets)
    - Can send transactions by creating and signing them; has no code
  - **Contract Accounts**
    - Controlled by their contract code; activates on every message received
    - Can read/write to internal storage, send messages, create contracts
- Every account has four fields:
  - nonce – transaction counter (prevents replay)
  - balance – amount of Wei ( $1 \text{ ETH} = 10^{18} \text{ Wei}$ )
  - storageRoot – hash of persistent key/value storage
  - codeHash – hash of EVM bytecode (empty for EOAs)

# Ethereum Account Types

I

## Externally Owned Account

(controlled by private key)

### Private Key

Signs transactions  
Derives address

**nonce** 7 TX count

**balance** 2.5 ETH Wei amount

**storageRoot** EMPTY —

**codeHash** Keccak256("") —

## Contract Account

(controlled by its own code)

### EVM Bytecode

Executes on every  
message/TX received  
PUSH MSTORE LT

### Persistent Storage

slot 0 → 0x42f...  
slot 1 → 0xab3...  
slot 2 → 0x000...

**nonce** 1 Created

**balance** 100 ETH Hold/send

**storageRoot** 0xa3f2e8... Merkle root

**codeHash** 0x7b9ef1... Keccak256

TX / Message Call

## Both account types share 20-byte address space (160 bits)

EOA = Keccak-256( publicKey ) [rightmost 20 bytes]

Contract = Keccak-256( RLP(sender, nonce) ) [rightmost 20 bytes]

EOAs are controlled by private keys; Contract Accounts are controlled by code

# Transactions and Messages

- Transaction: a signed data package from an EOA containing:
  - recipient – target address
  - signature – identifies the sender
  - value – amount of Ether to transfer
  - data – optional payload (e.g., function call + arguments)
  - gasLimit – max computational steps allowed
  - gasPrice – fee per computational step
- Messages: virtual objects sent between contracts (internal calls)
  - Like transactions but produced by contracts via CALL opcode
  - Never serialized; exist only in the EVM execution environment
- Gas allocation applies to the entire call chain:
  - If A sends TX with 1000 gas, B uses 600, calls C which uses 300, B has 100 left

# The Gas Mechanism

- Problem: Turing-complete code can loop forever – how to prevent DoS?
- Solution: every computational step costs gas
  - An attacker must pay proportionally for computation, bandwidth, and storage
- How it works:
  - Sender pre-pays  $\text{gasLimit} \times \text{gasPrice}$  in Ether
  - Each opcode consumes a defined amount of gas
  - Unused gas is refunded after execution
- Gas costs reflect resource consumption:
  - ADD = 3 gas, MUL = 5 gas, SSTORE = 20,000 gas (new slot)
  - Transaction data: 4 gas/zero byte, 16 gas/nonzero byte
- If execution runs out of gas:
  - All state changes revert, but gas is NOT refunded (prevents spam)

# The Gas Mechanism

## 1. User Submits TX

gasLimit = 100,000

gasPrice = 20 Gwei

Pre-pay:

$100,000 \times 20 \text{ Gwei}$

= 0.002 ETH upfront

## 2. EVM Executes

ADD → 3 gas

MUL → 5 gas

SLOAD → 2,100

SSTORE → 20,000

CALL → varies

Total: 65,000 gas

## 3a. SUCCESS ✓

State changes COMMIT

Refund unused gas:

$35,000 \times 20 \text{ Gwei}$

= 0.0007 ETH back

## 3b. OUT OF GAS ✗

ALL state changes

REVERT to original

Gas NOT refunded

(prevents spam)

## Why Gas Matters

- Turing-complete code loops forever
- Gas ensures every computation has cost
- Attacker pays for CPU, BW, storage
- Halting problem solved economically

## Selected Gas Costs (EIP-2929+)

**TX base:** 21,000 gas

**ADD:** 3 gas

**SSTORE:** 22,100 gas

**MUL:** 5 gas

**CREATE:** 32,000 gas

**SLOAD:** 2,100 gas

**CALL:** 2,600 gas

**BALANCE:** 2,600 gas

*Gas prevents infinite loops and DoS attacks by metering every computation*



# The Ethereum Virtual Machine

A Turing-complete, stack-based execution environment  
running identically on every node

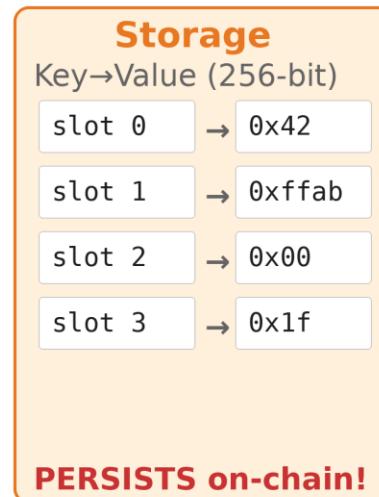
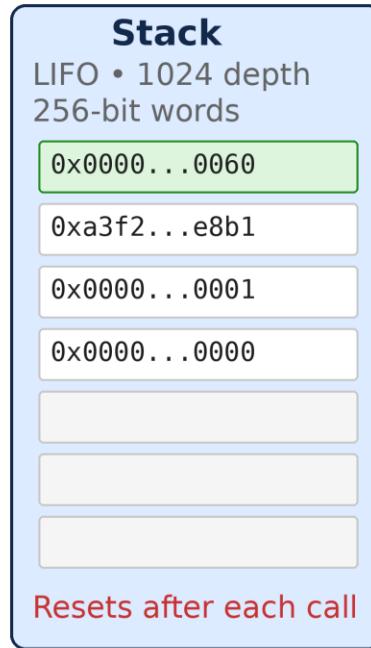
# EVM Architecture



- Stack-based virtual machine with 1024-item depth, 256-bit words
- Three data storage areas:
  - Stack – LIFO container; push/pop 256-bit values; resets after execution
  - Memory – expandable byte array; resets after execution
  - Storage – persistent key/value store (256-bit → 256-bit); survives between calls
- Execution model:
  - tuple (block\_state, transaction, message, code, memory, stack, pc, gas)
- Code = series of bytes; each byte is an opcode (ADD, SUB, SSTORE, CALL, ...)
  - Program counter increments through bytecode; STOP or RETURN halts
- Deterministic: given the same input, every node produces the same output
  - All nodes execute the same code – this IS the world computer

# EVM: Stack, Memory, Storage

I



Data Areas

PC→	0x00 STOP	0x01 ADD	0x02 MUL	0x10 LT
0x35	CALLDATALOAD	0x52 MSTORE	0x54 SLOAD	0x55 SSTORE
0x56	JUMP	0x57 JUMPI	0x60 PUSH1	0xF1 CALL
0xF3	RETURN	0xFD REVERT	0xF0 CREATE	0xFF SELFDESTRUCT

## Gas Counter

Remaining: 21,347 gas

Each opcode costs gas  
Out-of-gas = REVERT

## Execution Context

msg.sender – who called  
msg.value – ETH sent  
msg.data – calldata  
block.number, timestamp  
tx.origin – original EOA  
address(this) – contract

Execution Engine

The EVM is a stack-based machine with three distinct data areas

# Smart Contracts and Solidity

I

- Smart contracts:
  - Reactive agents, executing code when “poked” by a transaction or message
  - Have direct control over their own ETH balance and storage
- Solidity:
  - Primary high-level language for smart contracts; compiles to EVM
  - High-level, statically typed, object-oriented; influenced by C++, JavaScript, Python
- Development workflow:
  - Write Solidity → Compile to bytecode → Deploy → Interact
  - Tools: Remix (web IDE), Hardhat, Foundry
- Key properties:
  - Immutable once deployed (code cannot be changed)
  - Composable – contracts can call other contracts
- Other smart contract languages: Vyper (Python-like), Yul (low-level)

# ERC-20: The Token Standard

- A token is just a mapping: address → balance
  - Subtract X from A, give X to B (if A has at least X)
- ERC-20 (2015) standardizes this into 6 functions + 2 events:
  - balanceOf(owner) → uint256
  - transfer(to, amount) — direct send
  - approve(spender, amount) — grant spending rights
  - transferFrom(from, to, amount) — spend on behalf
  - totalSupply() and allowance(owner, spender)
  - Events: Transfer(from, to, value), Approval(owner, spender, value)
- Why it matters:
  - Any ERC-20 token works with every wallet, DEX, and dApp
  - USDC, USDT, UNI, LINK, DAI — all ERC-20
  - Composability: one standard → entire DeFi ecosystem

## SimpleToken.sol — Minimal ERC-20 Implementation

```
// SPDX-License-Identifier: MIT
pragma solidity ^0.8.20;

contract SimpleToken {
    string public name = "MyToken";
    string public symbol = "MTK";
    uint8 public decimals = 18;
    uint256 public totalSupply;

    mapping(address => uint256) public balanceOf;
    mapping(address => mapping(address => uint256))
        public allowance;

    event Transfer(address indexed from,
                  address indexed to, uint256 val);
    event Approval(address indexed owner,
                  address indexed spender, uint256);

    constructor(uint256 _supply) {
        totalSupply = _supply * 10**decimals;
        balanceOf[msg.sender] = totalSupply;
    }

    function transfer(address to, uint256 amt)
        public returns (bool) {
        require(balanceOf[msg.sender] >= amt);
        balanceOf[msg.sender] -= amt;
        balanceOf[to] += amt;
        emit Transfer(msg.sender, to, amt);
        return true;
    }
}
```

**State variables:**

name, symbol, decimals  
are token metadata

**The core state:**

mapping(addr → bal)

**transfer() is the heart of ERC-20:**

check → subtract → add  
→ emit event

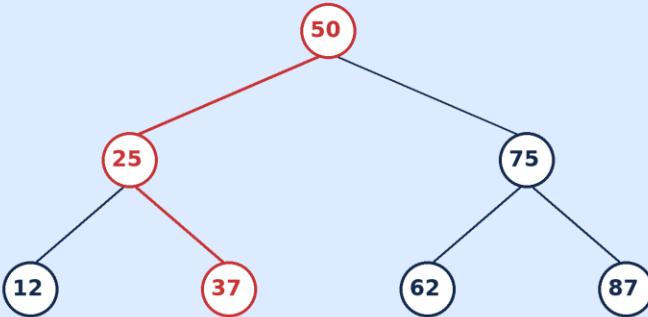
*A minimal but complete ERC-20 token — the building block of DeFi*

# Tree vs. Trie

## Binary Search Tree

Compare keys at each node

$O(\log n)$  lookup



### Find 37:

50 → go left

25 → go right

**37 → found!**

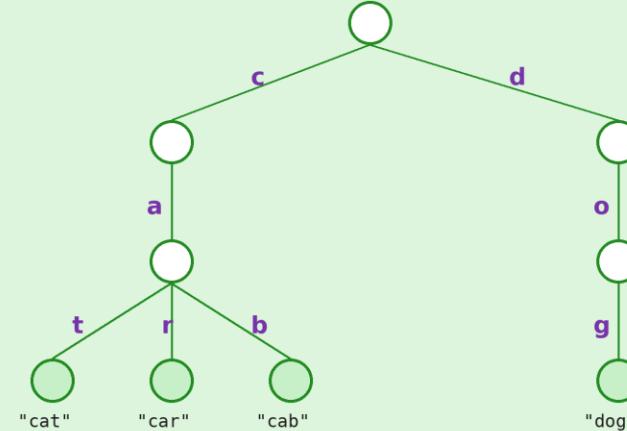
### Properties:

- Full key stored in each node
- Compare & branch left/right
- Good for ranges, bad for strings
- $O(\log n)$  — depends on # entries
- Non-deterministic (insert order matters)

## Trie (Prefix Tree)

Path from root encodes key

$O(\text{key length})$  lookup



### Find "car":

root → c → a → r → found!

Follow path edges, no comparisons

**3 steps = length of key**

### Properties:

- Key encoded in the path (edges)
- Shared prefixes = shared paths
- Perfect for address lookups
- $O(\text{key length})$  — independent of # entries
- **Deterministic (same keys → same trie)**

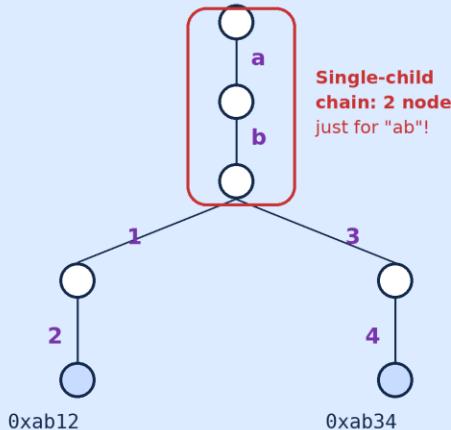
Ethereum needs deterministic lookups by address — tries provide exactly this

# Patricia Trie: Path Compression

I

## Standard Trie (wasteful)

Single-child chains waste space



Single-child  
chain: 2 nodes  
just for "ab"!

**7 nodes for just 2 keys**

Each nibble = one node, one edge

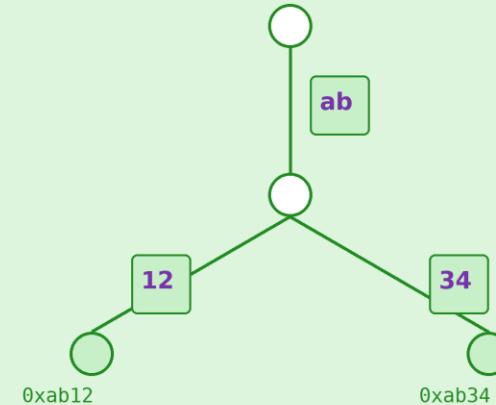
→ Wastes nodes where there is no branching

### Why Ethereum needs this:

- Keys = Keccak-256 hashes, 40 nibbles (hex digits)
- Max path depth: 40 edges
- Patricia compresses non-branching segments
- Each internal node stores hash of its subtree
- Merkle proof:  $O(\log n)$  hashes to verify
- **Result: Modified Merkle Patricia Trie**
- = trie + path compression + Merkle hashing

## Patricia Trie (compressed)

Collapse single-child chains



**4 nodes for 2 keys!**

"ab" collapsed into single edge

Branch only where keys diverge

### Key insight:

Edges store multi-character strings,  
not single characters. This is what makes  
it "Patricia" (Practical Algorithm To  
Retrieve Information Coded In Alphanumeric)

Patricia = Practical Algorithm To Retrieve Information Coded In Alphanumeric

# State Storage: Modified Merkle Patricia Trie

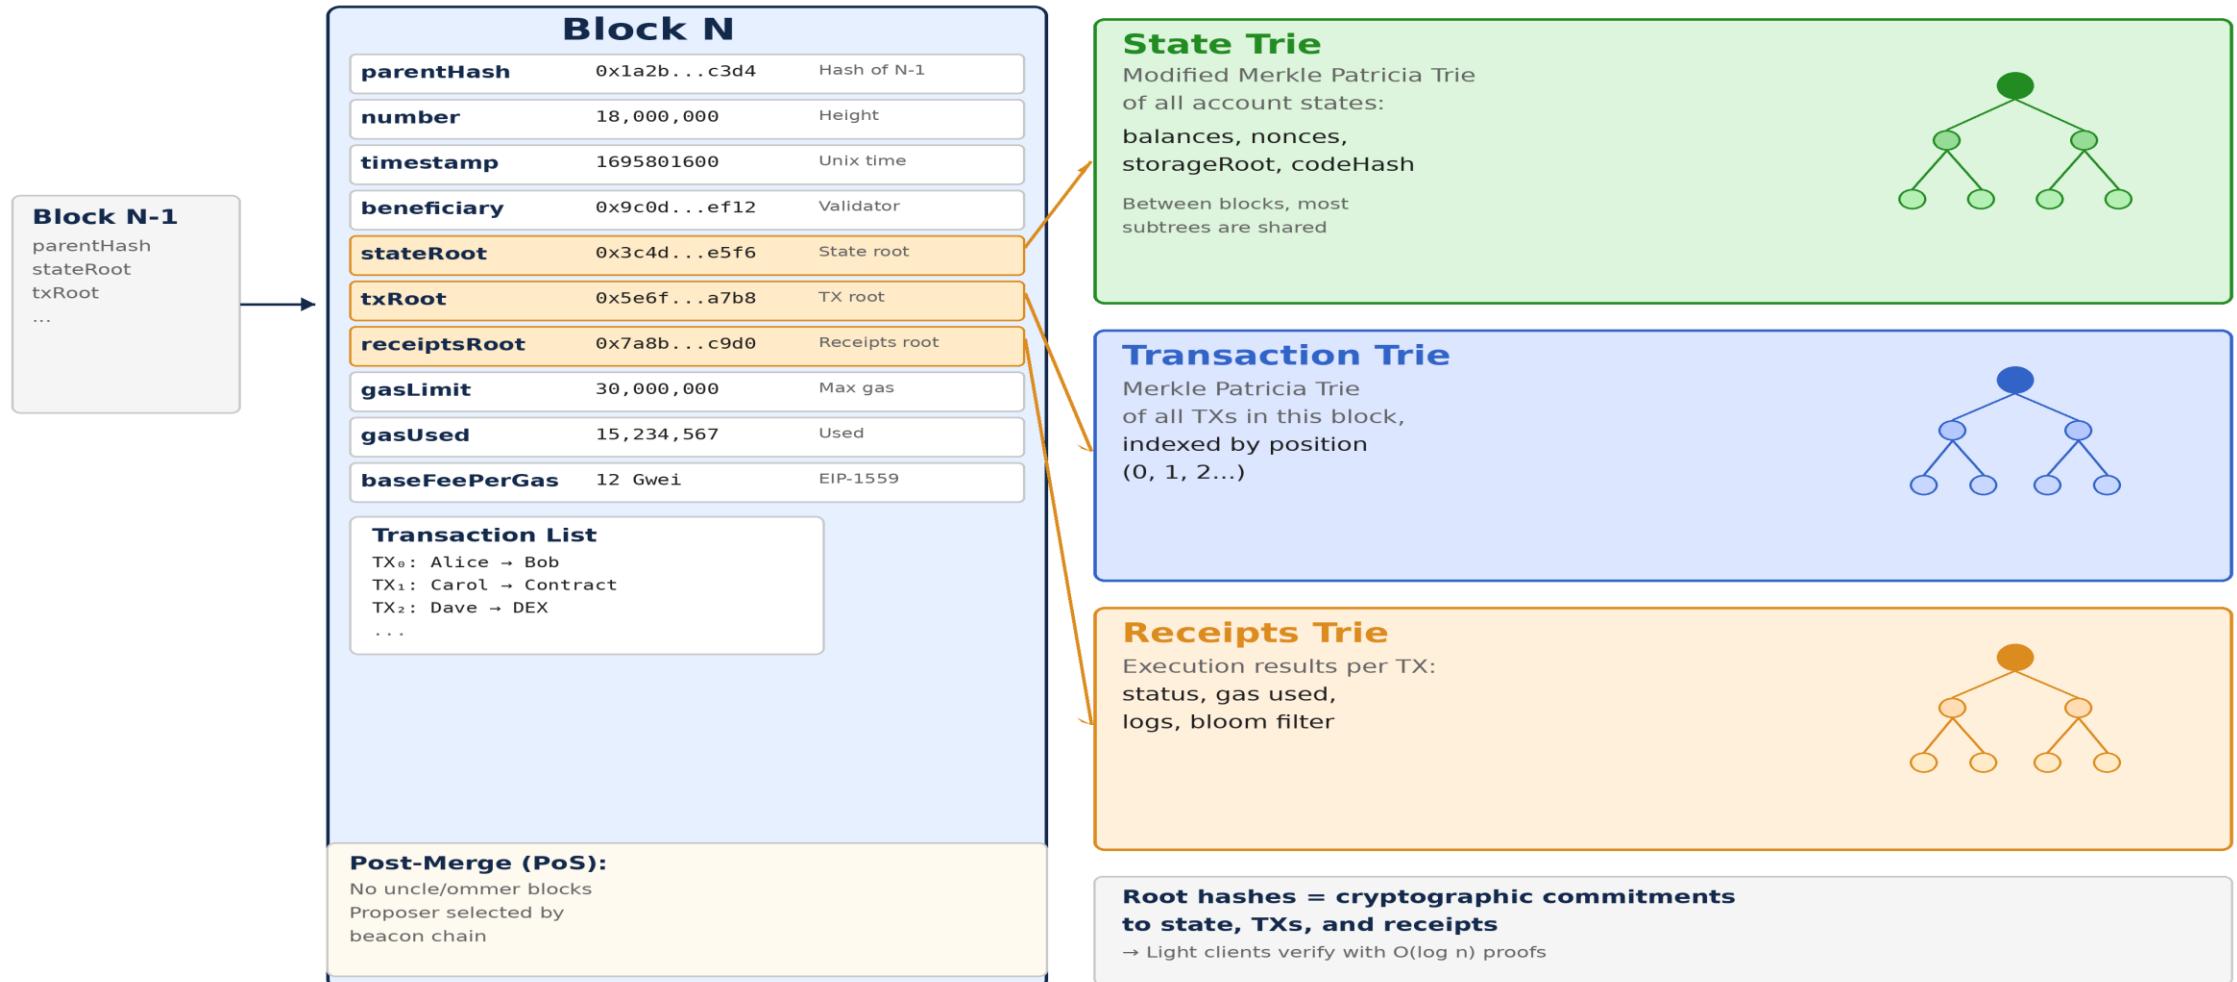
- Ethereum state is stored in a Patricia Trie (not just a Merkle tree)
  - Root hash in block header = cryptographic commitment to entire world state
- Why not a simple Merkle tree?
  - Need to insert and delete accounts, not just verify membership
  - Need efficient key-value lookups, not just ordered data
- Three tries in every block header:
  - State trie – all account states
  - Receipts trie – execution results for each transaction
  - Transaction trie – all transactions in the block
- Efficiency: between adjacent blocks, most of the trie is identical

# Block Structure and Validation

- Ethereum blocks store more than Bitcoin blocks:
  - Parent hash, timestamp, block number, difficulty
  - State root, transaction root, receipts root (three tries!)
  - Gas limit, gas used, beneficiary (validator/miner address)
- Block validation algorithm (simplified):
  - 1. Verify previous block exists and is valid
  - 2. Check timestamp constraints
  - 3. For each TX:  $S[i+1] = \text{APPLY}(S[i], \text{TX}[i])$ ; check total gas  $\leq \text{GASLIMIT}$
  - 4. Verify final state root matches header
- GHOST Protocol (Greedy Heaviest Observed Subtree):
  - Includes “uncle” (stale) blocks in chain weight calculation
  - Uncle blocks receive 87.5% of base reward; nephew gets 12.5%
  - Reduces centralization bias from network latency

# Ethereum Block Structure

I



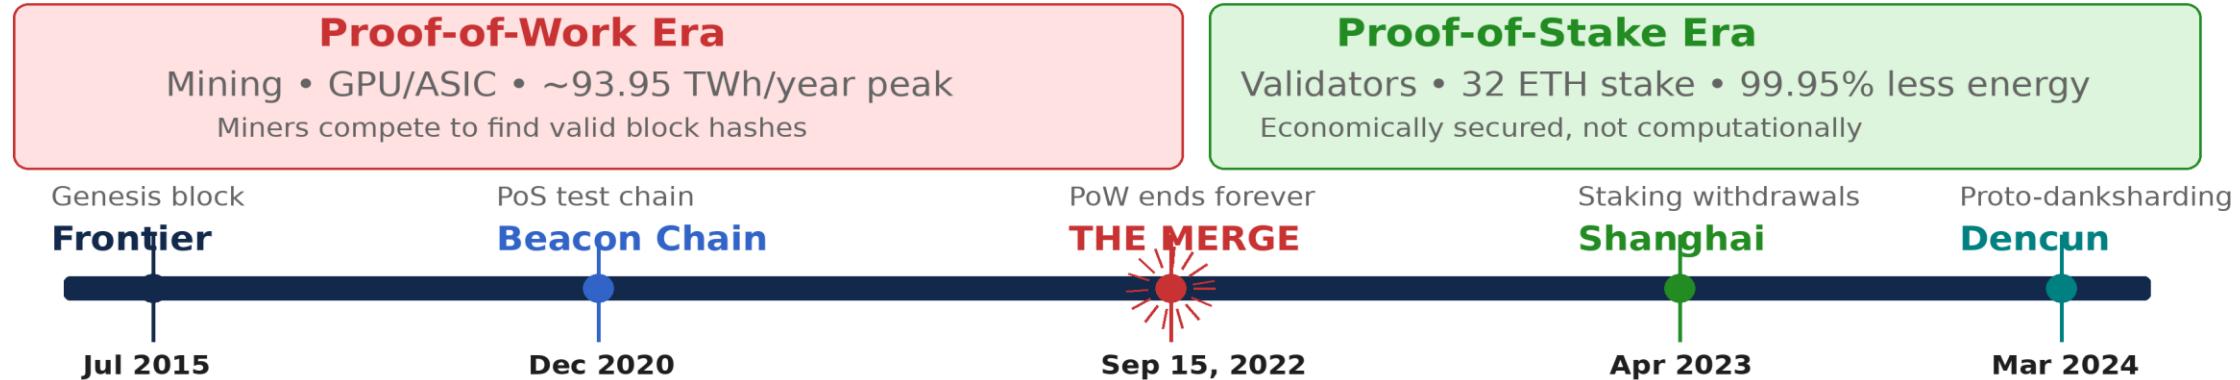
Three Merkle tries provide cryptographic commitments to state, transactions, and receipts

# The Merge: From Proof-of-Work to Proof-of-Stake

I

- Ran Proof-of-Work from 2015 until The Merge (Sept 2022)
- Beacon Chain launched December 1, 2020 – a parallel PoS chain
  - Ran in parallel; validators staked 32 ETH; tested PoS consensus independently
- The Merge: September 15, 2022
  - Execution layer (transactions) merged with consensus layer (Beacon Chain)
  - Mining permanently deprecated; PoS validators now produce blocks
  - Triggered by reaching Terminal Total Difficulty (difficulty bomb)
- Impact:
  - ~99.95% reduction in energy consumption
  - Block time stabilized at exactly 12 seconds
  - Total block reward income (USD) fell by ~97%
- No transaction history was lost – first PoS block attached to last PoW block

# The Merge: PoW to PoS Timeline



## Impact of The Merge

<b>Energy:</b>	<b>~99.95%</b>	$\sim 93.95 \rightarrow \sim 0.01 \text{ TWh/yr}$
<b>Block time:</b>	<b>Exactly 12 sec</b>	Steady, no PoW variance
<b>Reward:</b>	<b>-97%</b>	TX fee revenue increased
<b>Validators:</b>	<b>500,000+</b>	Lido, Coinbase, Rocket Pool
<b>TX history:</b>	<b>Fully preserved</b>	No data lost

## How The Merge Worked

1. Beacon Chain launched Dec 2020
2. Ran parallel ~21 months testing
3. Validators staked 32 ETH
4. Terminal Total Difficulty reached
5. Execution layer swapped to PoS
6. Mining permanently discontinued

# Proof-of-Stake Mechanics

- Validators replace miners; must stake 32 ETH to participate
  - Over 500,000 validators active post-merge
- Every 12 seconds (one slot), the protocol:
  - Randomly selects a block proposer from the validator set
  - Assigns a committee of attesters (~1/32 of all validators)
  - Attesters vote on the proposed block's validity
- Slashing: penalties for malicious behavior
  - Double-signing or contradictory attestations → lose portion of staked ETH
  - Inactivity leak: offline validators gradually lose stake
- Finality: Casper FFG (Friendly Finality Gadget)
  - Checkpoints every 32 slots (1 epoch ~6.4 min); finalized after 2 epochs
- Shanghai upgrade (April 2023) finally enabled staking withdrawals

# Ethereum Applications: DeFi, NFTs, DAOs

- DeFi (Decentralized Finance)
  - Automated Market Makers (Uniswap), lending (Aave, Compound)
  - Stablecoins (DAI, USDC) – hedging contracts from the whitepaper, realized
  - Over \$100B total value locked at peak
- NFTs (Non-Fungible Tokens – ERC-721)
  - Unique digital ownership: art, collectibles, gaming items, real-world assets
  - Each token has a unique ID; ownership tracked on-chain
- DAOs (Decentralized Autonomous Organizations)
  - On-chain governance with token-weighted voting
  - Treasury management without traditional legal structures
  - The original whitepaper vision of “decentralized autonomous organizations”
- Other applications:
  - Identity systems (ENS), prediction markets, decentralized storage, layer-2 rollups

# Ethereum Milestones

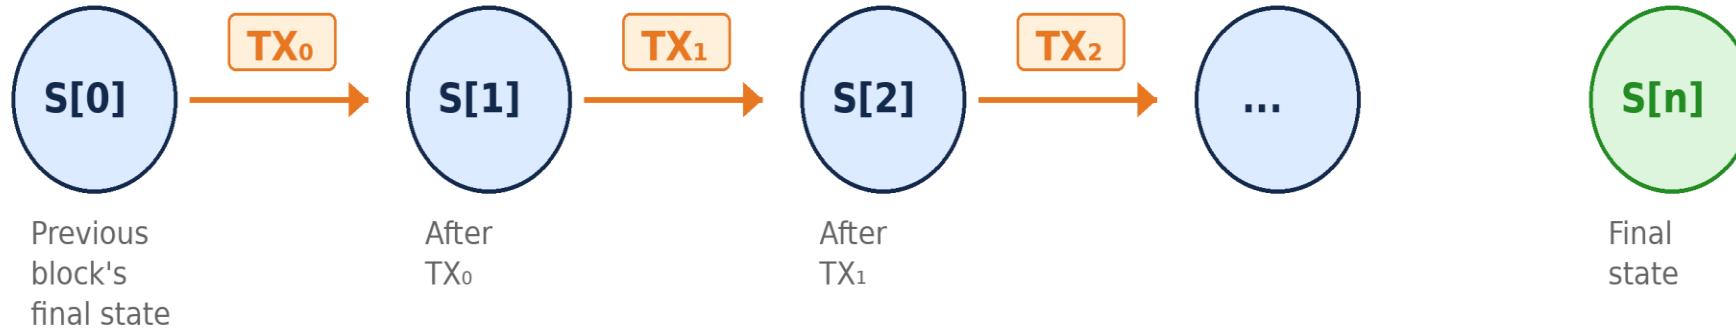


- Nov 2013 – Vitalik Buterin publishes the Ethereum Whitepaper
- Jul 2014 – Public crowdsale raises ~\$18M (31,500 BTC)
- Jul 30, 2015 – Frontier launch (genesis block)
- Jun 2016 – The DAO hack (\$60M); hard fork creates ETH/ETC split
- Dec 2017 – CryptoKitties congests the network; scaling debate intensifies
- Jan 2018 – ETH price peaks near \$1,400
- Aug 2021 – EIP-1559 (London): base fee burn makes ETH deflationary
- Sep 15, 2022 – The Merge: PoW → PoS
- Apr 2023 – Shanghai: staking withdrawals enabled
- 2024–2025 – Dencun (proto-danksharding) and Pectra upgrades

- The Yellow Paper defines the EVM – but in informal mathematics
  - Ambiguities and inconsistencies discovered over the years
- KEVM: a complete, executable formal semantics of the EVM in the K Framework
  - Developed at UIUC and Runtime Verification
  - The “Jello Paper” – [jellopaper.org](http://jellopaper.org) – a readable presentation of the semantics
- What makes KEVM special:
  - Complete – passes the entire official Ethereum test suite
  - Executable – can run EVM programs directly
  - Formal – enables mathematical proofs about smart contract correctness
- Practical impact:
  - Formal verification of smart contracts (e.g., ERC-20 implementations)
  - Found real bugs in the Yellow Paper specification

# Block Processing: Sequential State Transitions

I



**$S[i+1] = \text{APPLY}( S[i] , \text{TX}[i] )$**

$S[0]$  = previous block's final state

If ANY returns ERROR → block INVALID

$S[n]$  = final state → Merkle root in header

## Block Validation:

1. Previous block exists and is valid
2. Timestamp in valid range
3. Total gas  $\leq$  GASLIMIT

**4. MerkleRoot( $S[n]$ ) == stateRoot ✓**

*Each transaction is applied sequentially; the final state root must match the header*

# Ethereum vs. Bitcoin: How It All Fits Together

- State model: UTXO (Bitcoin) vs. Accounts (Ethereum)
  - Stateless vs. stateful; simple transfers vs. complex interactions
- Scripting: limited stack-based (Bitcoin) vs. Turing-complete EVM (Ethereum)
- Consensus: PoW (Bitcoin) vs. PoS (Ethereum, since 2022)
- Block time: ~10 min (Bitcoin) vs. ~12 sec (Ethereum)
- Data in blocks: transactions only (Bitcoin) vs. transactions + state root (Ethereum)
- Supply: capped at 21M BTC vs. no hard cap, but deflationary since EIP-1559
- Purpose: peer-to-peer electronic cash vs. general-purpose world computer
  - Digital gold vs. decentralized application platform