



THE GRAINGER COLLEGE OF ENGINEERING

CS 521

Technological Foundations of Blockchain and Cryptocurrency

Grigore Rosu

Topic 4 – Ethereum

I ILLINOIS

Bitcoin's Limitations – Why Something More?

I

- Lack of Turing-completeness
 - Script language supports no loops; simulating them requires duplicating code
- Value-blindness
 - UTXO is all-or-nothing; no fine-grained control over withdrawal amounts
- Lack of state
 - UTXO are either spent or unspent; no multi-stage contracts or persistent variables
- Blockchain-blindness
 - Scripts cannot access block data (nonce, timestamp, previous hash)



Ethereum

A next-generation smart contract
and decentralized application platform

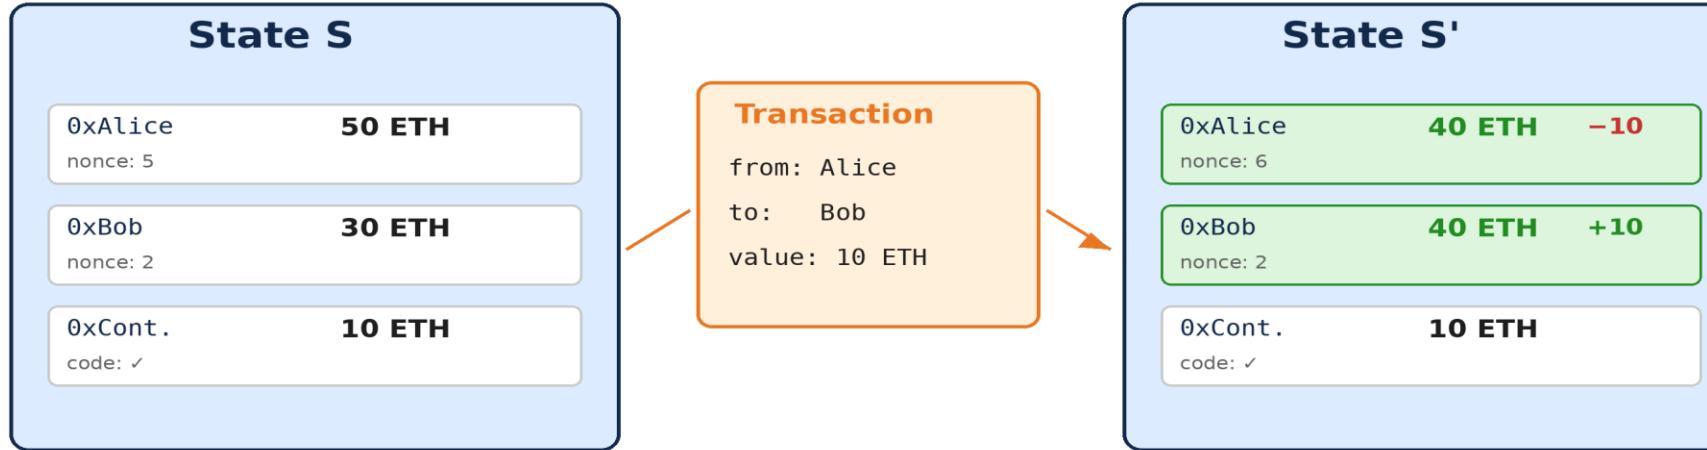
Ethereum: A Programmable Blockchain

- Proposed by Vitalik Buterin in late 2013; whitepaper published 2014
- Goal: a blockchain with a built-in Turing-complete programming language
- Anyone can write smart contracts and decentralized applications (dApps)
 - Users define their own logic in a few lines of code
- Create arbitrary rules for ownership, transaction formats, and state transitions
- Not just digital cash – a general-purpose decentralized world computer
 - Tokens, DeFi, NFTs, DAOs, identity, governance, and more
- Launched July 30, 2015 (Frontier release)

Ethereum as a State Transition System

- State: mapping from addresses (160-bit) to account states
 - Every account has: nonce, balance, storage root, code hash
- State transition function: $\text{APPLY}(S, \text{TX}) \rightarrow S'$ or ERROR
 - Validates signature, deducts gas fees, runs code, refunds unused gas
- Transactions morph the state incrementally, block by block
- Key difference from Bitcoin:
 - Bitcoin = UTXO model (stateless); Ethereum = account model (stateful)
- Ethereum blocks contain both the transaction list AND the most recent state (root hash)
- State stored in a Modified Merkle Patricia Trie
 - Allows efficient state lookups and proofs without storing full history

Ethereum State Transition Function



APPLY(S , TX) → S' or ERROR

Yellow Paper: $\sigma(t+1) \equiv Y(\sigma(t), T)$

APPLY({ Alice: 50, Bob: 30 }, "send 70")

→ **ERROR (insufficient balance: 50 < 70)**

From the Ethereum Whitepaper: APPLY(S, TX) produces a new state S'

Ethereum Accounts

- Two types of accounts sharing the same 20-byte address space:
 - **Externally Owned Accounts (EOAs)**
 - Controlled by private keys (humans/wallets)
 - Can send transactions by creating and signing them; has no code
 - **Contract Accounts**
 - Controlled by their contract code; activates on every message received
 - Can read/write to internal storage, send messages, create contracts
- Every account has four fields:
 - nonce – transaction counter (prevents replay)
 - balance – amount of Wei ($1 \text{ ETH} = 10^{18} \text{ Wei}$)
 - storageRoot – hash of persistent key/value storage
 - codeHash – hash of EVM bytecode (empty for EOAs)

Ethereum Account Types

Externally Owned Account

(controlled by private key)

Private Key

Signs transactions
Derives address

nonce	7	TX count
--------------	---	----------

balance	2.5 ETH	Wei amount
----------------	---------	------------

storageRoot	EMPTY	—
--------------------	-------	---

codeHash	Keccak256("")	—
-----------------	---------------	---

TX / Message Call

Contract Account

(controlled by its own code)

EVM Bytecode

Executes on every message/TX received
PUSH 0x60 MSTORE
CALLDATASIZE LT

Persistent Storage

slot 0 → 0x42f...
slot 1 → 0xab3...
slot 2 → 0x000...
Survives between TXs

Both account types share the same 20-byte address space (160 bits)

EOA address = Keccak-256(publicKey) [rightmost 20 bytes]

Contract address = Keccak-256(RLP(sender, nonce)) [rightmost 20 bytes]

EOAs are controlled by private keys; Contract Accounts are controlled by code

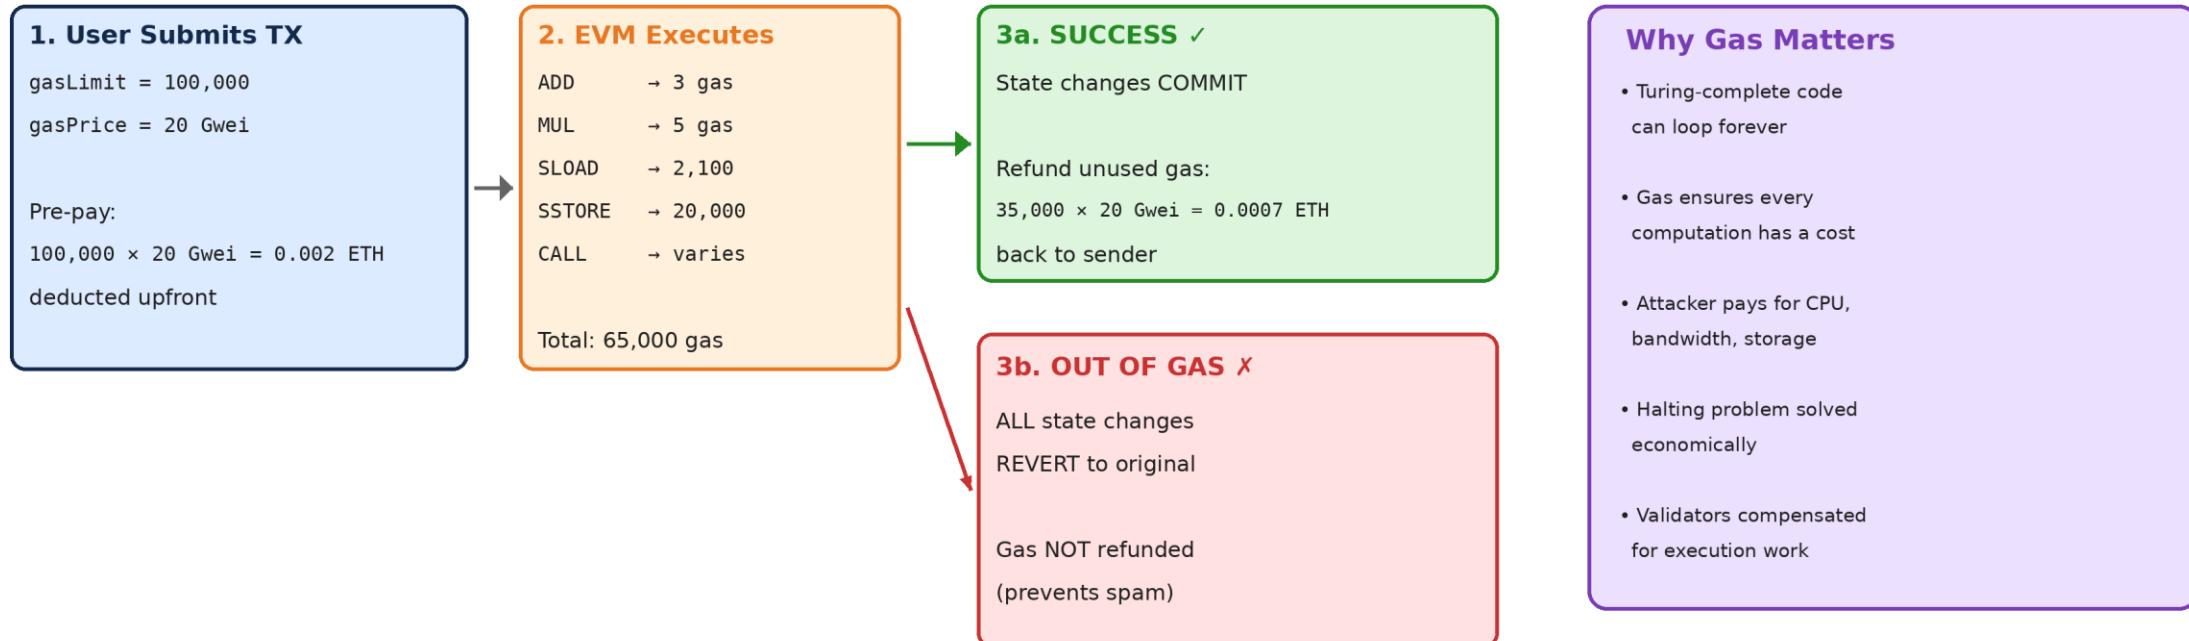
Transactions and Messages

- Transaction: a signed data package from an EOA containing:
 - recipient – target address
 - signature – identifies the sender
 - value – amount of Ether to transfer
 - data – optional payload (e.g., function call + arguments)
 - gasLimit – max computational steps allowed
 - gasPrice – fee per computational step
- Messages: virtual objects sent between contracts (internal calls)
 - Like transactions but produced by contracts via CALL opcode
 - Never serialized; exist only in the EVM execution environment
- Gas allocation applies to the entire call chain:
 - If A sends TX with 1000 gas, B uses 600, calls C which uses 300, B has 100 left

The Gas Mechanism

- Problem: Turing-complete code can loop forever – how to prevent DoS?
- Solution: every computational step costs gas
 - An attacker must pay proportionally for computation, bandwidth, and storage
- How it works:
 - Sender pre-pays $\text{gasLimit} \times \text{gasPrice}$ in Ether
 - Each opcode consumes a defined amount of gas
 - Unused gas is refunded after execution
- Gas costs reflect resource consumption:
 - ADD = 3 gas, MUL = 5 gas, SSTORE = 20,000 gas (new slot)
 - Transaction data: 4 gas/zero byte, 16 gas/nonzero byte
- If execution runs out of gas:
 - All state changes revert, but gas is NOT refunded (prevents spam)

The Gas Mechanism



Selected Gas Costs (EIP-2929+)							
TX base fee:	21,000 gas	SSTORE (new):	22,100 gas	CREATE:	32,000 gas	CALL (cold):	2,600 gas
ADD / SUB:	3 gas	MUL / DIV:	5 gas	SLOAD (cold):	2,100 gas	SLOAD (warm):	100 gas
BALANCE:	2,600 gas	LOGO:	375 gas	CALLDATACOPY:	3+3/word gas	EXTCODESIZE:	2,600 gas

Gas prevents infinite loops and DoS attacks by metering every computation

The Ethereum Virtual Machine

A Turing-complete, stack-based execution environment
running identically on every node

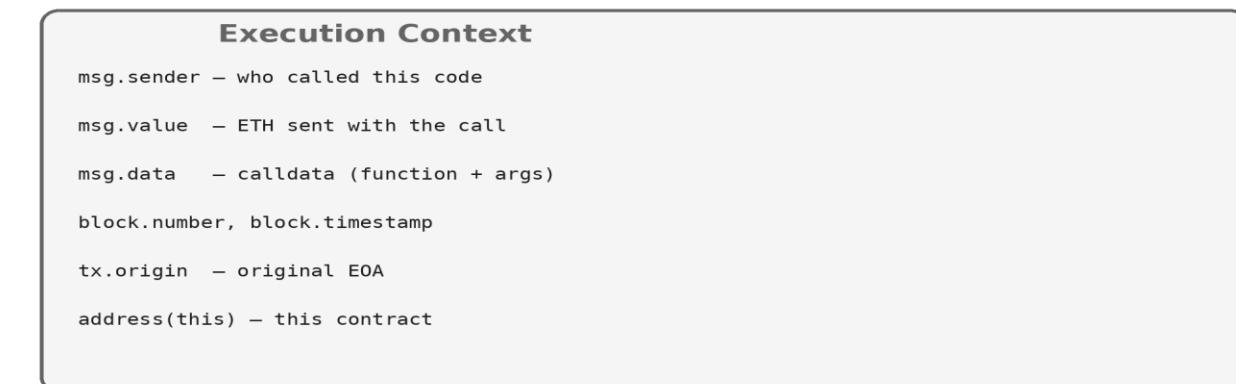
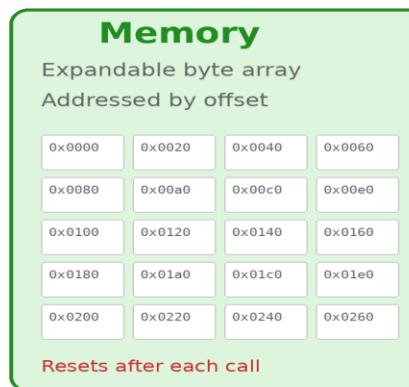
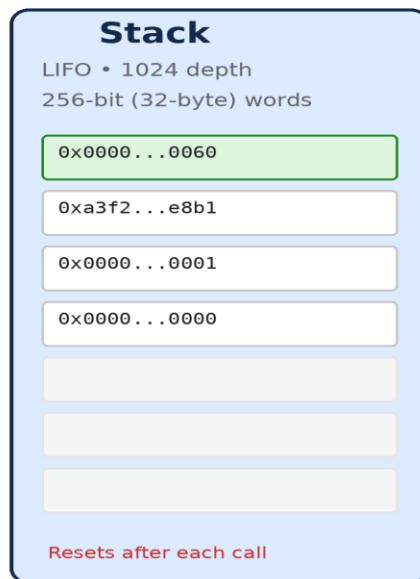
EVM Architecture



- Stack-based virtual machine with 1024-item depth, 256-bit words
- Three data storage areas:
 - Stack – LIFO container; push/pop 256-bit values; resets after execution
 - Memory – expandable byte array; resets after execution
 - Storage – persistent key/value store (256-bit → 256-bit); survives between calls
- Execution model: tuple (block_state, transaction, message, code, memory, stack, pc, gas)
- Code = series of bytes; each byte is an opcode (ADD, SUB, SSTORE, CALL, ...)
 - Program counter increments through bytecode; STOP or RETURN halts
- Deterministic: given the same input, every node produces the same output
 - All nodes execute the same code – this IS the world computer

EVM: Stack, Memory, Storage

The logo of the University of Illinois, consisting of a large blue letter 'I' with a white outline, set against an orange background.



The EVM is a stack-based machine with three distinct data areas

Smart Contracts and Solidity

- Smart contracts: autonomous agents living in the Ethereum execution environment
 - Execute code when “poked” by a transaction or message
 - Have direct control over their own Ether balance and storage
- Solidity: the primary high-level language for writing smart contracts
 - High-level, statically typed, object-oriented; influenced by C++, JavaScript, Python
 - Compiles to EVM bytecode
- Development workflow:
 - Write contract in Solidity → Compile to bytecode → Deploy via transaction → Interact
 - Tools: Remix (web IDE), Hardhat, Foundry
- Key properties:
 - Immutable once deployed (code cannot be changed)
 - Composable – contracts can call other contracts
- Other smart contract languages: Vyper (Python-like), Yul (low-level)

A Concrete Smart Contract: Token System

- All a token system fundamentally is: a database with one operation
 - Subtract X units from A and give X units to B, if A has at least X
- The basic logic (from Ethereum Whitepaper):
 - def send(to, value):
 - if self.storage[msg.sender] >= value:
 - self.storage[msg.sender] -= value
 - self.storage[to] += value
- ERC-20 standard formalizes the interface:
 - totalSupply(), balanceOf(address), transfer(to, value)
 - approve(spender, value), transferFrom(from, to, value)
 - Any token following ERC-20 works with every wallet, DEX, and dApp
- This elegance is why Ethereum became the platform for thousands of tokens

State Storage: Modified Merkle Patricia Trie

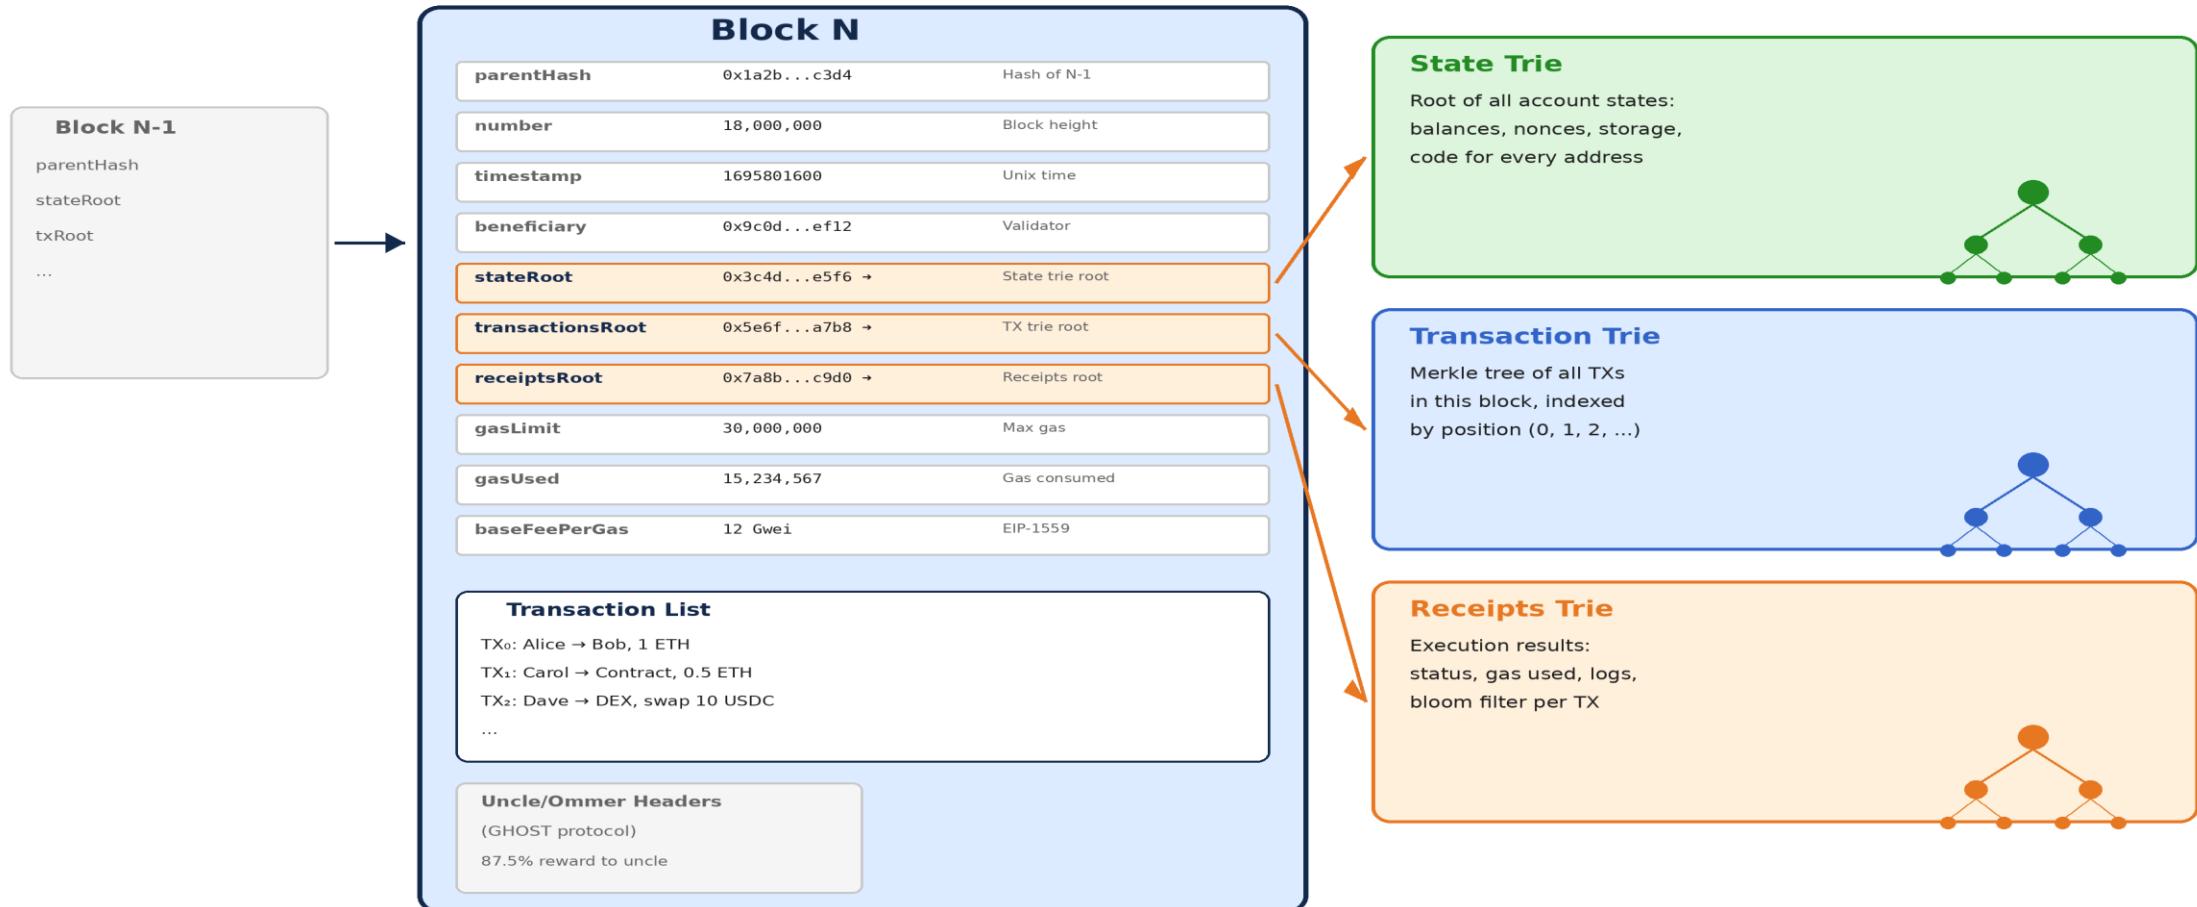
- Ethereum state is stored in a Patricia Trie (not just a Merkle tree)
 - Root hash in block header = cryptographic commitment to entire world state
- Why not a simple Merkle tree?
 - Need to insert and delete accounts, not just verify membership
 - Need efficient key-value lookups, not just ordered data
- Patricia Trie combines:
 - Radix trie (compact prefix tree) for efficient lookups by address
 - Merkle hashing for tamper-evidence and light client proofs
- Three tries in every block header:
 - State trie – all account states
 - Storage trie – each contract's persistent data
 - Transaction trie – all transactions in the block
- Efficiency: between adjacent blocks, most of the trie is identical

Block Structure and Validation

- Ethereum blocks store more than Bitcoin blocks:
 - Parent hash, timestamp, block number, difficulty
 - State root, transaction root, receipts root (three tries!)
 - Gas limit, gas used, beneficiary (validator/miner address)
- Block validation algorithm (simplified):
 - 1. Verify previous block exists and is valid
 - 2. Check timestamp constraints
 - 3. For each TX: $S[i+1] = \text{APPLY}(S[i], \text{TX}[i])$; check total gas $\leq \text{GASLIMIT}$
 - 4. Verify final state root matches header
- GHOST Protocol (Greedy Heaviest Observed Subtree):
 - Includes “uncle” (stale) blocks in chain weight calculation
 - Uncle blocks receive 87.5% of base reward; nephew gets 12.5%
 - Reduces centralization bias from network latency

Ethereum Block Structure

I



Three Merkle tries provide cryptographic commitments to state, transactions, and receipts

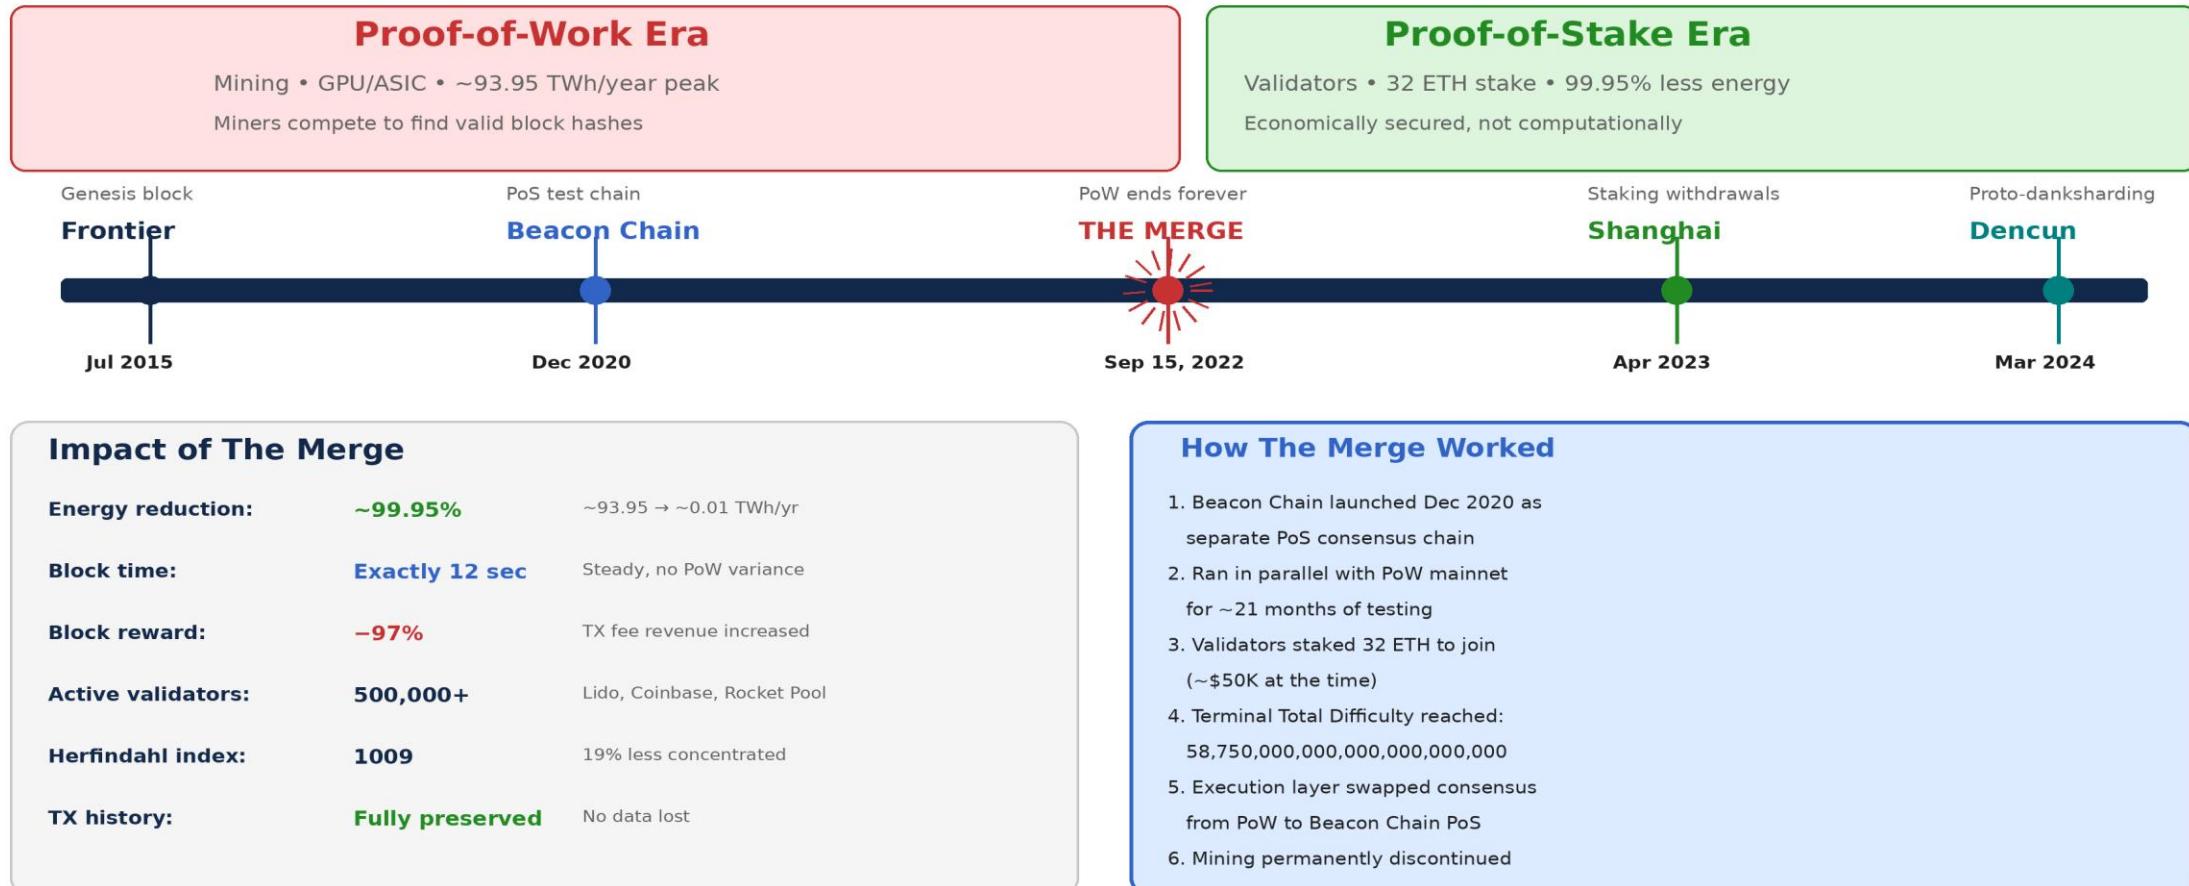
The Merge: From Proof-of-Work to Proof-of-Stake



- Ethereum ran Proof-of-Work from launch (July 2015) until September 2022
- Beacon Chain launched December 1, 2020 – a parallel PoS chain
 - Ran in parallel; validators staked 32 ETH; tested PoS consensus independently
- The Merge: September 15, 2022
 - Execution layer (transactions) merged with consensus layer (Beacon Chain)
 - Mining permanently deprecated; PoS validators now produce blocks
 - Triggered by reaching Terminal Total Difficulty (difficulty bomb)
- Impact:
 - ~99.95% reduction in energy consumption
 - Block time stabilized at exactly 12 seconds
 - Total block reward income (USD) fell by ~97%
- No transaction history was lost – first PoS block attached to last PoW block

The Merge: PoW to PoS Timeline

I



September 15, 2022: Ethereum completes the transition to Proof-of-Stake

Proof-of-Stake Mechanics

- Validators replace miners; must stake 32 ETH to participate
 - Over 500,000 validators active post-merge
- Every 12 seconds (one slot), the protocol:
 - Randomly selects a block proposer from the validator set
 - Assigns a committee of attesters (~1/32 of all validators)
 - Attesters vote on the proposed block's validity
- Slashing: penalties for malicious behavior
 - Double-signing or contradictory attestations → lose portion of staked ETH
 - Inactivity leak: offline validators gradually lose stake
- Finality: Casper FFG (Friendly Finality Gadget)
 - Checkpoints every 32 slots (1 epoch ~6.4 min); finalized after 2 epochs
- Shanghai upgrade (April 2023) finally enabled staking withdrawals

Ethereum Applications: DeFi, NFTs, DAOs

- DeFi (Decentralized Finance)
 - Automated Market Makers (Uniswap), lending (Aave, Compound)
 - Stablecoins (DAI, USDC) – hedging contracts from the whitepaper, realized
 - Over \$100B total value locked at peak
- NFTs (Non-Fungible Tokens – ERC-721)
 - Unique digital ownership: art, collectibles, gaming items, real-world assets
 - Each token has a unique ID; ownership tracked on-chain
- DAOs (Decentralized Autonomous Organizations)
 - On-chain governance with token-weighted voting
 - Treasury management without traditional legal structures
 - The original whitepaper vision of “decentralized autonomous organizations”
- Other applications:
 - Identity systems (ENS), prediction markets, decentralized storage, layer-2 rollups

Ethereum Milestones



- Nov 2013 – Vitalik Buterin publishes the Ethereum Whitepaper
- Jul 2014 – Public crowdsale raises ~\$18M (31,500 BTC)
- Jul 30, 2015 – Frontier launch (genesis block)
- Jun 2016 – The DAO hack (\$60M); hard fork creates ETH/ETC split
- Dec 2017 – CryptoKitties congests the network; scaling debate intensifies
- Jan 2018 – ETH price peaks near \$1,400
- Aug 2021 – EIP-1559 (London): base fee burn makes ETH deflationary
- Sep 15, 2022 – The Merge: PoW → PoS
- Apr 2023 – Shanghai: staking withdrawals enabled
- 2024–2025 – Dencun (proto-danksharding) and Pectra upgrades

KEVM: Formal Semantics of EVM (The Jello Paper)

I

- The Yellow Paper defines the EVM – but in informal mathematics
 - Ambiguities and inconsistencies discovered over the years
- KEVM: a complete, executable formal semantics of the EVM in the K Framework
 - Developed at UIUC and Runtime Verification
 - The “Jello Paper” – jellopaper.org – a readable presentation of the semantics
- What makes KEVM special:
 - Complete – passes the entire official Ethereum test suite
 - Executable – can run EVM programs directly
 - Formal – enables mathematical proofs about smart contract correctness
- Practical impact:
 - Formal verification of smart contracts (e.g., ERC-20 implementations)
 - Found real bugs in the Yellow Paper specification
 - Used by Runtime Verification for auditing high-value DeFi protocols
- KEVM paper published at CSF 2018 (IEEE Computer Security Foundations)

Block Processing: Sequential State Transitions

I



$$S[i+1] = \text{APPLY}(S[i] , TX[i])$$

$S[0]$ = state at end of previous block

If ANY returns ERROR → entire block is INVALID

$S[n]$ = final state → Merkle root stored in block header

Block Validation:

1. Previous block exists and is valid
2. Timestamp in valid range
3. Total gas \leq GASLIMIT
4. **MerkleRoot($S[n]$) == stateRoot in header ✓**

Each transaction is applied sequentially; the final state root must match the header

Ethereum vs. Bitcoin: How It All Fits Together

I

- State model: UTXO (Bitcoin) vs. Accounts (Ethereum)
 - Stateless vs. stateful; simple transfers vs. complex interactions
- Scripting: limited stack-based (Bitcoin) vs. Turing-complete EVM (Ethereum)
- Consensus: PoW (Bitcoin) vs. PoS (Ethereum, since 2022)
- Block time: ~10 min (Bitcoin) vs. ~12 sec (Ethereum)
- Data in blocks: transactions only (Bitcoin) vs. transactions + state root (Ethereum)
- Supply: capped at 21M BTC vs. no hard cap, but deflationary since EIP-1559
- Purpose: peer-to-peer electronic cash vs. general-purpose world computer
 - Digital gold vs. decentralized application platform
- Both are decentralized, permissionless, and secured by cryptographic primitives